

Transactions & Concurrency Control II

Alvin Cheung
Aditya Parameswaran
Reading: R & G Chapter 16-17



Transaction Implementations



- Many available!
 - Targets different workloads
- We will focus on **lock-based** implementations
 - Others: use multiple versions of data and “optimistically” let transactions move forward
 - Abort when conflicts are detected
 - Some names to know/look up:
 - Optimistic Concurrency Control
 - Timestamp-Ordered Multiversion Concurrency Control
 - We will not study these schemes in this lecture

“Lock” data??

- Not by any crypto or hardware enforcement
 - There are no adversaries here ... this is all within the DBMS
- Recall locks / semaphores from 61c
 - These are synchronization primitives
 - Locking / unlocking has costs
- We lock by simple convention within the DBMS:
 - Each *data element* has a unique **lock**
 - Each transaction must first **acquire** the lock before reading/writing that element
 - If the lock is taken by another transaction, then wait
 - The transaction must **release** the lock(s) at some point
- Different *lock protocols / schemes* differ by:
 - When to lock / unlock each data element
 - What data element to lock
 - What happens when a txn waits for a lock



What are “data elements”?



Major differences between vendors:

- Lock on the entire database
 - SQLite
- Lock on individual records
 - SQL Server, DB2, etc
- Will see tradeoffs later on

Actions on Locks



$\text{Lock}_i(A) / L_i(A)$ = transaction T_i acquires lock for element A

$\text{Unlock}_i(A) / U_i(A)$ = transaction T_i releases lock for element A

Let's see this in action...

A Non-S Serializable Schedule



T1

READ(A)
A := A+100
WRITE(A)

READ(B)
B := B+100
WRITE(B)

T2

READ(A)
A := A*2
WRITE(A)
READ(B)
B := B*2
WRITE(B)

Example

T1

$L_1(A)$; READ(A)

$A := A + 100$

WRITE(A); $U_1(A)$; $L_1(B)$

READ(B)

$B := B + 100$

WRITE(B); $U_1(B)$;

T2

$L_2(A)$; READ(A)

$A := A * 2$

WRITE(A); $U_2(A)$;

$L_2(B)$; BLOCKED...

...GRANTED; READ(B)

$B := B * 2$

WRITE(B); $U_2(B)$;



Using locks has ensured a conflict-serializable schedule

But...



T1

$L_1(A)$; READ(A)
A := A+100
WRITE(A); $U_1(A)$;

T2

$L_2(A)$; READ(A)
A := A*2
WRITE(A); $U_2(A)$;
 $L_2(B)$; READ(B)
B := B*2
WRITE(B); $U_2(B)$;



$L_1(B)$; READ(B)
B := B+100
WRITE(B); $U_1(B)$;

Locks did not enforce conflict-serializability!!! What's wrong ?

Two Phase Locking (2PL)



The 2PL rule:

In every transaction, all lock requests must precede all unlock requests

Example: 2PL transactions



T1

$L_1(A); L_1(B);$ READ(A)
 $A := A + 100$
WRITE(A); $U_1(A)$

READ(B)
 $B := B + 100$
WRITE(B); $U_1(B)$

Now it is conflict-serializable

T2

$L_2(A);$ READ(A)
 $A := A * 2$
WRITE(A);
 $L_2(B);$ BLOCKED...

...GRANTED; READ(B)
 $B := B * 2$
WRITE(B); $U_2(A); U_2(B)$

Two Phase Locking (2PL)



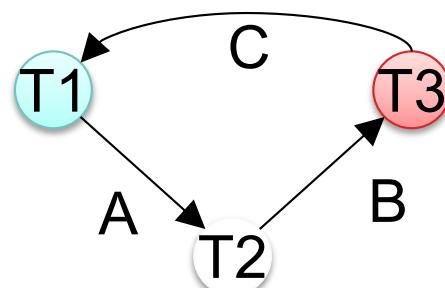
Theorem: 2PL ensures conflict serializability

Two Phase Locking (2PL)



Theorem: 2PL ensures conflict serializability

Proof. Suppose not: then there exists a cycle in the dependence graph.



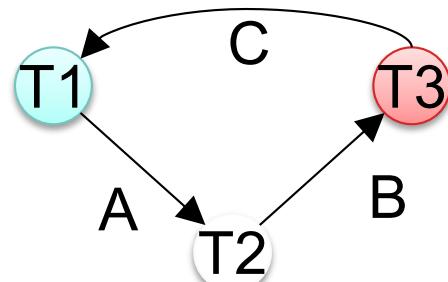
Two Phase Locking (2PL)



Theorem: 2PL ensures conflict serializability

Proof. Suppose not: then there exists a cycle in the dependence graph.

Then there is the following temporal cycle in the schedule:

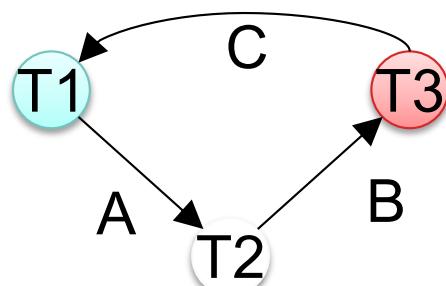


Two Phase Locking (2PL)



Theorem: 2PL ensures conflict serializability

Proof. Suppose not: then there exists a cycle in the dependence graph.



Then there is the following temporal cycle in the schedule:

$U_1(A) \rightarrow L_2(A)$ why?

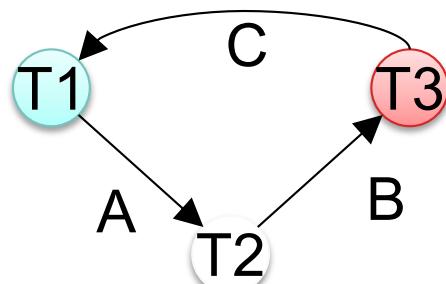
$U_1(A)$ happened strictly before $L_2(A)$

Two Phase Locking (2PL)



Theorem: 2PL ensures conflict serializability

Proof. Suppose not: then there exists a cycle in the dependence graph.



Then there is the following temporal cycle in the schedule:

$U_1(A) \rightarrow L_2(A)$

$L_2(A) \rightarrow U_2(B)$

why?

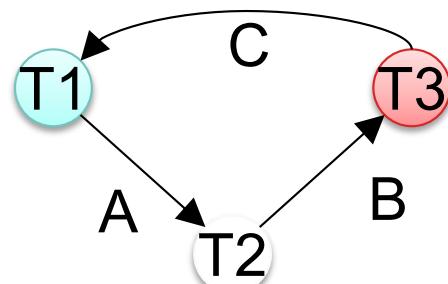
$L_2(A)$ happened strictly before $U_1(A)$

Two Phase Locking (2PL)



Theorem: 2PL ensures conflict serializability

Proof. Suppose not: then there exists a cycle in the dependence graph.



Then there is the following temporal cycle in the schedule:

$U_1(A) \rightarrow L_2(A)$

$L_2(A) \rightarrow U_2(B)$

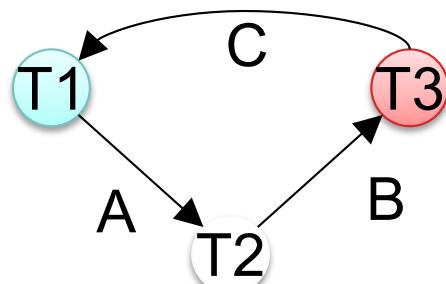
$U_2(B) \rightarrow L_3(B)$ why?

Two Phase Locking (2PL)



Theorem: 2PL ensures conflict serializability

Proof. Suppose not: then there exists a cycle in the dependence graph.



Then there is the following temporal cycle in the schedule:

$U_1(A) \rightarrow L_2(A)$

$L_2(A) \rightarrow U_2(B)$

$U_2(B) \rightarrow L_3(B)$

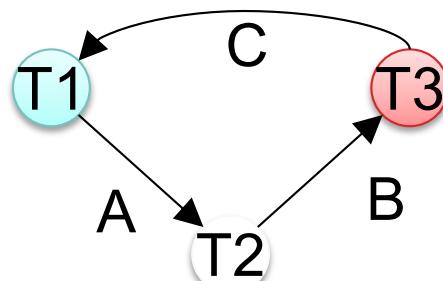
.....etc.....

Two Phase Locking (2PL)



Theorem: 2PL ensures conflict serializability

Proof. Suppose not: then there exists a cycle in the dependence graph.



Then there is the following temporal cycle in the schedule:

$U_1(A) \rightarrow L_2(A)$
 $L_2(A) \rightarrow U_2(B)$
 $U_2(B) \rightarrow L_3(B)$
 $L_3(B) \rightarrow U_3(C)$
 $U_3(C) \rightarrow L_1(C)$
 $L_1(C) \rightarrow U_1(A)$

Cycle in time:
Contradiction

A New Problem: Non-recoverable Schedule



T1

$L_1(A); L_1(B);$ READ(A)

$A := A + 100$

WRITE(A); $U_1(A)$

READ(B)

$B := B + 100$

WRITE(B); $U_1(B)$;

Rollback

T2

$L_2(A);$ READ(A)

$A := A * 2$

WRITE(A);

$L_2(B);$ BLOCKED...

...GRANTED; READ(B)

$B := B * 2$

WRITE(B); $U_2(A); U_2(B);$

Commit

Aka cascading aborts

A New Problem: Non-recoverable Schedule



T1

$L_1(A); L_1(B);$ READ(A)
A := A+100
WRITE(A); $U_1(A)$

READ(B)
B := B+100
WRITE(B); $U_1(B)$

Rollback

Elements A, B written
by T1 are restored
to their original value.

T2

$L_2(A);$ READ(A)
A := A*2
WRITE(A);
 $L_2(B);$ BLOCKED...

...GRANTED; READ(B)
B := B*2
WRITE(B); $U_2(A); U_2(B)$;
Commit

A New Problem: Non-recoverable Schedule



T1

$L_1(A); L_1(B);$ READ(A)
A := A+100
WRITE(A); $U_1(A)$

READ(B)
B := B+100
WRITE(B); $U_1(B)$

Rollback

Elements A, B written
by T1 are restored
to their original value.

T2

$L_2(A);$ READ(A)
A := A*2
WRITE(A);
 $L_2(B);$ BLOCKED...

...GRANTED; READ(B)
B := B*2
WRITE(B); $U_2(A); U_2(B)$; Commit

Dirty reads of
A, B lead to
incorrect writes.

A New Problem: Non-recoverable Schedule



T1

$L_1(A); L_1(B); \text{READ}(A)$
 $A := A + 100$
 $\text{WRITE}(A); U_1(A)$

$\text{READ}(B)$
 $B := B + 100$
 $\text{WRITE}(B); U_1(B);$

Rollback

Elements A, B written
by T1 are restored
to their original value.

T2

$L_2(A); \text{READ}(A)$
 $A := A * 2$
 $\text{WRITE}(A);$
 $L_2(B); \text{BLOCKED...}$

...GRANTED; $\text{READ}(B)$
 $B := B * 2$
 $\text{WRITE}(B); U_2(A); U_2(B);$
Commit

Can no longer undo!

Dirty reads of
A, B lead to
incorrect writes.



Strict 2PL



The Strict 2PL rule:

All locks are held until commit/abort:

All unlocks are done together with commit/abort.

With strict 2PL, we will get schedules that are both conflict-serializable and recoverable

Strict 2PL



T1

$L_1(A)$; READ(A)

$A := A + 100$

WRITE(A);

$L_1(B)$; READ(B)

$B := B + 100$

WRITE(B);

Rollback & $U_1(A); U_1(B)$;

T2

$L_2(A)$; BLOCKED...

...GRANTED; READ(A)

$A := A * 2$

WRITE(A);

$L_2(B)$; READ(B)

$B := B * 2$

WRITE(B);

Commit & $U_2(A); U_2(B)$;

Strict 2PL



- Lock-based systems always use strict 2PL
- Easy to implement:
 - Before a transaction reads or writes an element A, insert an L(A)
 - When the transaction commits/aborts, then release all locks
- Ensures both conflict serializability and recoverability

Another problem: Deadlocks



- T_1 : $R(A), W(B)$
 - T_2 : $R(B), W(A)$
-
- T_1 holds the lock on A, waits for B
 - T_2 holds the lock on B, waits for A



This is a deadlock!

Deadlock Prevention

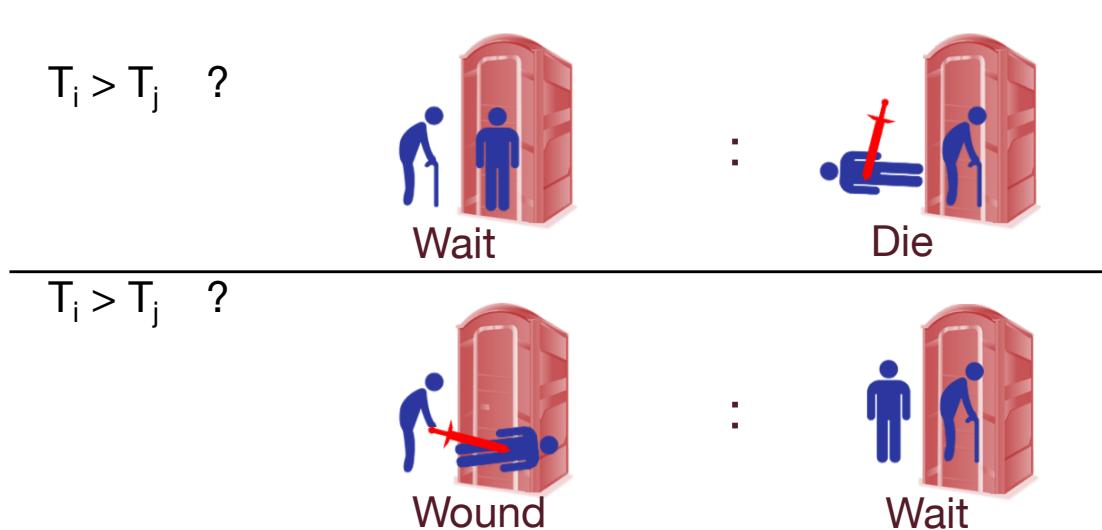


- **Common technique in operating systems**
- **Standard approach: resource ordering**
 - Screen < Network Card < Printer
- **Why is this problematic for transactions?**
 - What order would you impose?

Deadlock Avoidance



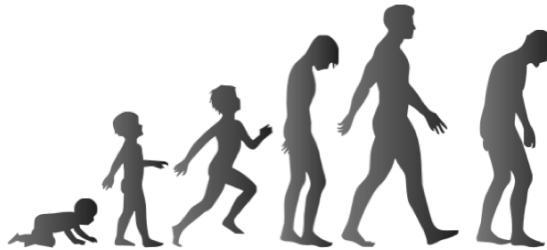
- **Assign priorities based on age: (now – start_time).**
- Say T_i wants a lock that T_j holds. Two possible policies:
 - **Wait-Die:** If T_i has higher priority, T_i waits for T_j ; else T_i aborts
 - **Wound-Wait:** If T_i has higher priority, T_j aborts; else T_i waits
 - Read each of these like a ternary operator (C/C++/java/javascript)



Deadlock Avoidance: Analysis



- Q: Why do these schemes guarantee no deadlocks?
 - Q: What do the previous images have in common?
- Important Detail: If a transaction re-starts, make sure it gets its original timestamp. Why?
- Note: other priority schemes make sense
 - E.g. measures of resource consumption, like #locks acquired



Deadlock Detection



- Create and maintain a **“waits-for” graph**
- Periodically check for cycles in a graph

Deadlock Detection, Part 2

Example:

T1:

T2:

T3:

T4:



T1

T2

T4

T3

Deadlock Detection, Part 3

Example:

T1: R(A)

T2:

T3:

T4:



T1

T2

T4

T3

Deadlock Detection, Part 4

Example:

T1: R(A) R(D)

T2:

T3:

T4:



T1

T2

T4

T3

Deadlock Detection, Part 5

Example:

T1: R(A) R(D)

T2: W(B)

T3:

T4:



T1

T2

T4

T3

Deadlock Detection, Part 6

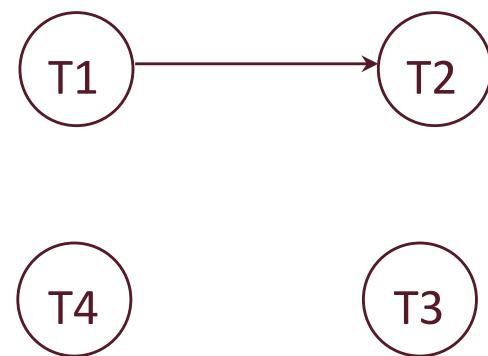
Example:

T1: R(A) R(D) R(B)

T2: W(B)

T3:

T4:



Deadlock Detection, Part 7

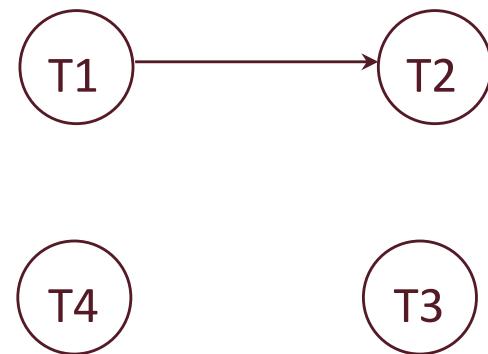
Example:

T1: R(A) R(D) R(B)

T2: W(B)

T3: R(D)

T4:



Deadlock Detection, Part 8

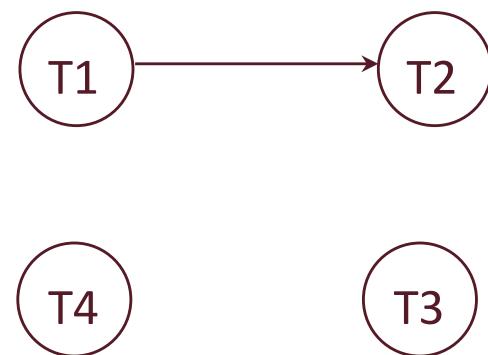
Example:

T1: R(A) R(D) R(B)

T2: W(B)

T3: R(D) R(C)

T4:



Deadlock Detection, Part 9

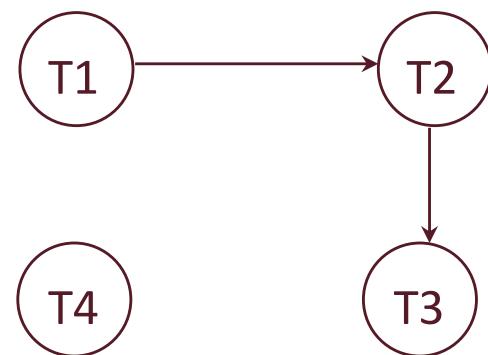
Example:

T1: R(A) R(D) R(B)

T2: W(B) W(C)

T3: R(D) R(C)

T4:



Deadlock Detection, Part 10

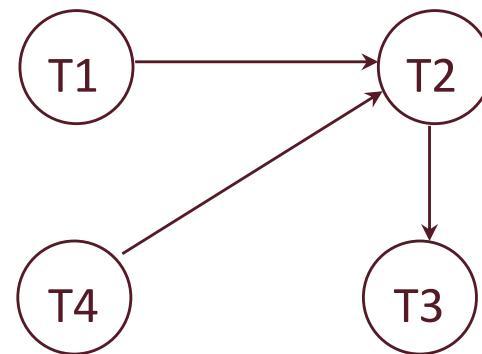
Example:

T1: R(A) R(D) R(B)

T2: W(B) W(C)

T3: R(D) R(C)

T4: W(B)



Deadlock Detection, Part 11

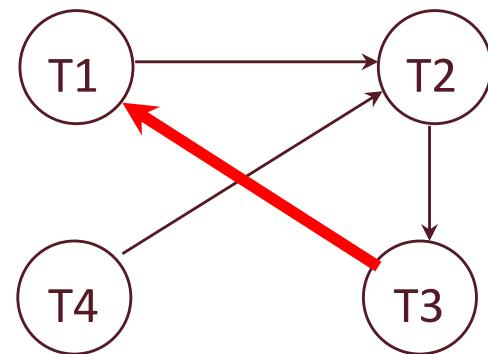
Example:

T1: R(A) R(D) R(B)

T2: W(B) W(C)

T3: R(D) R(C)

T4: W(A)
W(B)



Deadlock!



- T1, T2, T3 are deadlocked
 - Doing no good, and holding locks
- T4 still cruising
- In the background, run a deadlock detection algorithm
 - Periodically extract the waits-for graph
 - Find cycles
 - “Shoot” a transaction on the cycle
- Empirical fact
 - Most deadlock cycles are small (2-3 transactions)

Lock Modes



- **S** = shared lock (for READ)
- **X** = exclusive lock (for WRITE)
- Cannot get new locks after releasing any locks (strict 2PL)

Lock compatibility matrix:

	None	S	X
None			
S			
X			

Lock Modes



- **S** = shared lock (for READ)
- **X** = exclusive lock (for WRITE)
- Cannot get new locks after releasing any locks (strict 2PL)

Lock compatibility matrix:

	None	S	X
None	✓	✓	✓
S	✓	✓	✗
X	✓	✗	✗

Lock Management



- Lock and unlock requests handled by Lock Manager
- LM maintains a hashtable, keyed on names of objects being locked.
- LM keeps an entry for each currently held lock
- Entry contains
 - Granted set: Set of txns currently granted access to the lock
 - Lock mode: Type of lock held (shared or exclusive)
 - Wait Queue: Queue of lock requests

	Granted Set	Mode	Wait Queue
A	{T1, T2}	S	T3(X) ← T4(X)
B	{T6}	X	T5(X) ← T7(S)

Lock Management (continued)



- **When lock request arrives:**
 - Does any txn in Granted Set or Wait Queue want a conflicting lock?
 - If no, put the requester into “granted set” and let them proceed
 - If yes, put requester into wait queue (typically FIFO)
- **Lock upgrade:**
 - Txn with shared lock can request to upgrade to exclusive

	Granted Set	Mode	Wait Queue
A	{T1, T2}	S	T3(X) ← T4(X)
B	{T6}	X	T5(X) ← T7(S)

Lock Granularity



- **Fine granularity locking** (e.g., tuples)
 - High concurrency
 - High overhead in managing locks
 - E.g., SQL Server
- **Coarse grain locking** (e.g., tables, entire database)
 - Many false conflicts
 - Less overhead in managing locks
 - E.g., SQL Lite
- **Solution: lock escalation changes granularity as needed**

Lock Granularity, cont



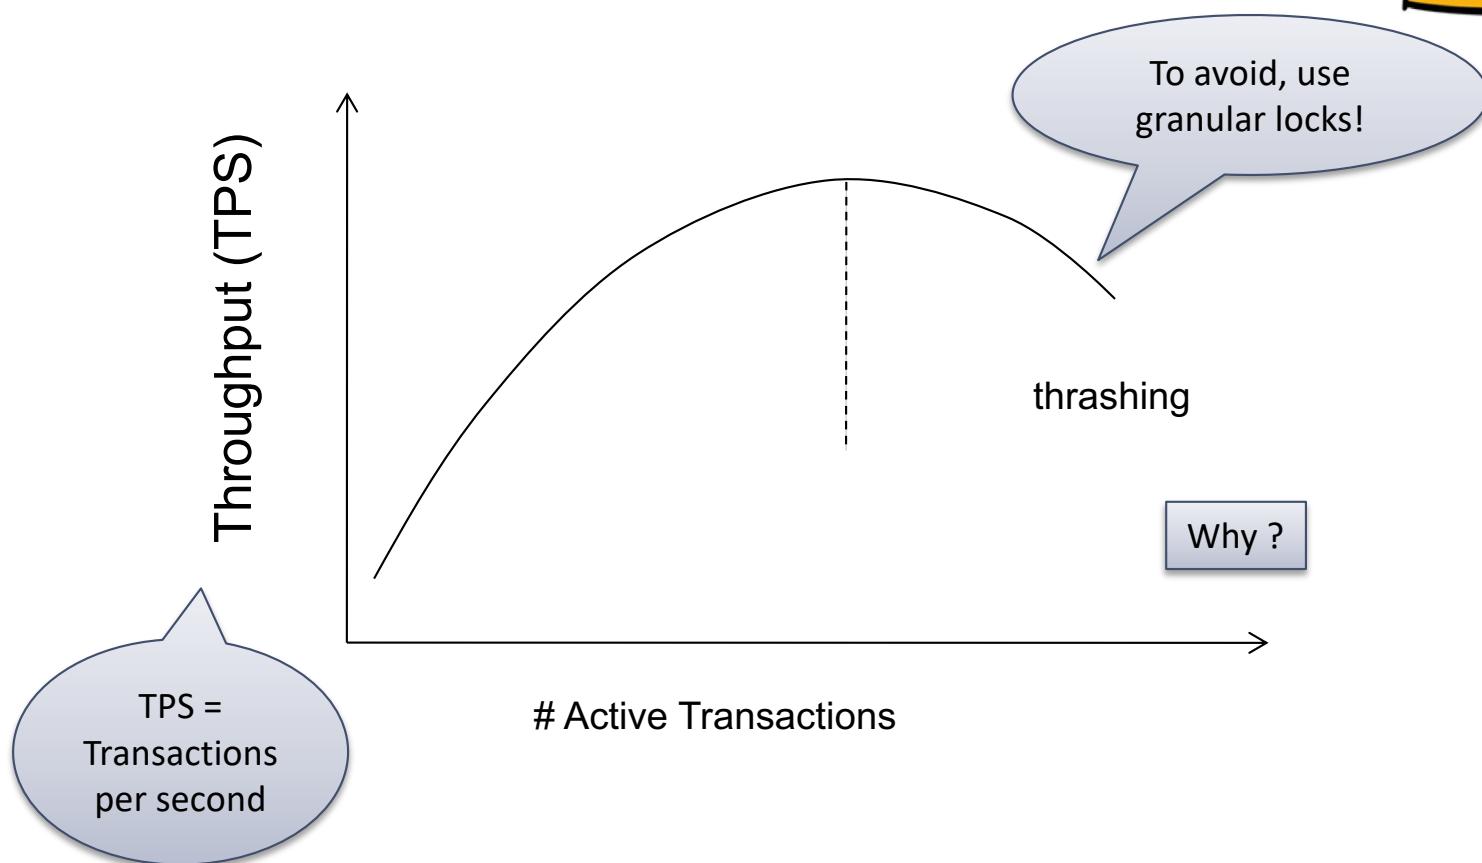
- Hard to decide what granularity to lock
 - Tuples vs Pages vs Tables?
- What is the tradeoff?
 - Fine-grained availability of resources would be nice (e.g. lock per tuple)
 - Small # of locks to manage would also be nice (e.g. lock per table)
 - Can't have both!
 - Or can we???

Multiple Locking Granularity



- **Shouldn't have to make same decision for all transactions!**
- Allow data items to be of various sizes
- Define a hierarchy of data granularities, small nested within large
 - Can be represented graphically as a tree.

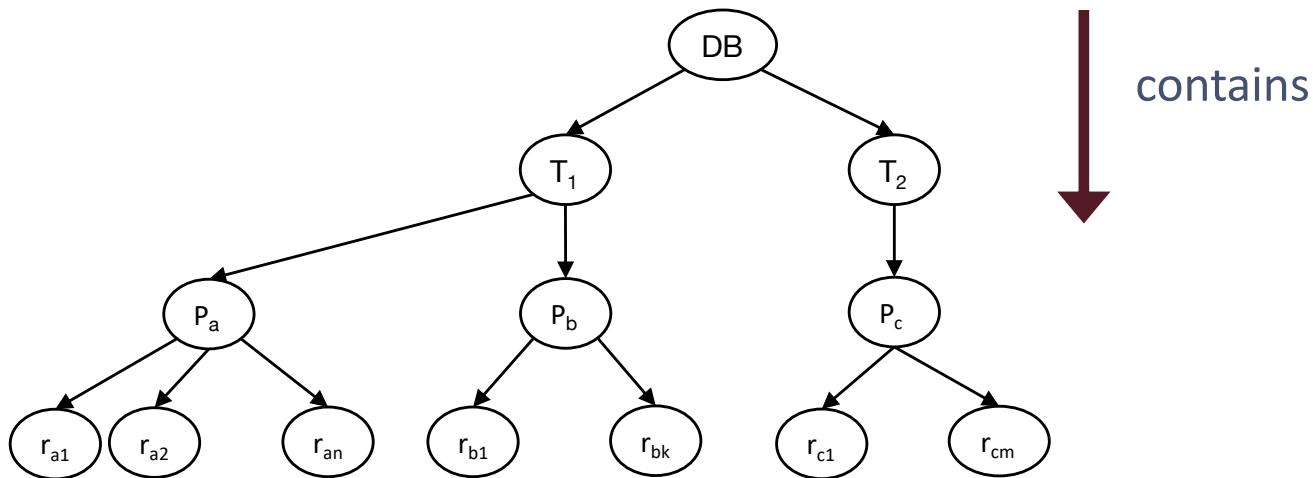
Lock Performance



Example of Granularity Hierarchy



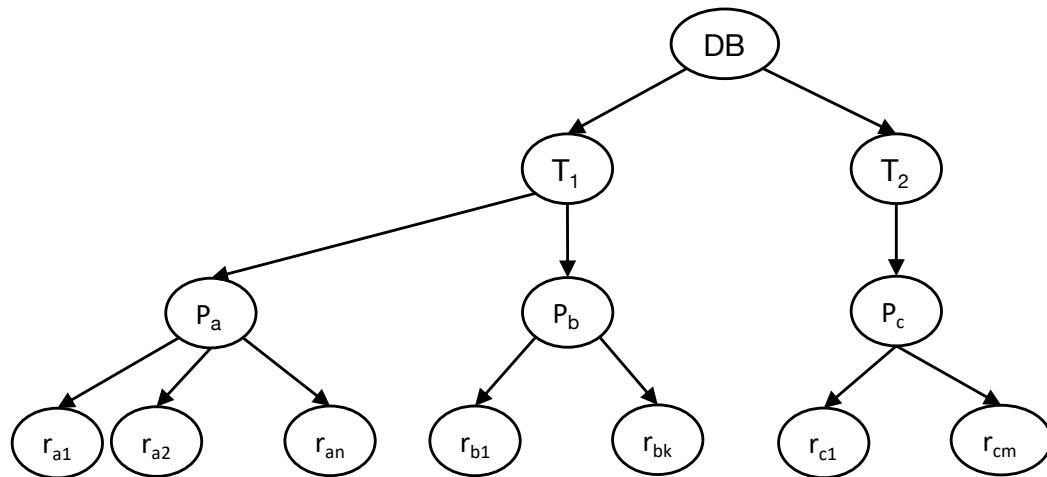
- Data “containers” can be viewed as nested.
- The levels, starting from the coarsest (top) level are
 - Database, Tables, Pages, Records
- When a transaction locks a node in the tree **explicitly**, it **implicitly** locks all the node’s descendants in the same mode.



Multiple Locking Granularity



- Granularity of locking (level in tree where locking is done):
 - **Fine granularity** (lower in tree): High concurrency, lots of locks (overhead)
 - **Coarse granularity** (higher in tree): Few locks (low overhead), lost concurrency
 - Lost potential concurrency if you don't need everything inside the coarse grain



Real-World Locking Granularities



Resource	Description
RID	A row identifier used to lock a single row within a heap.
KEY	A row lock within an index used to protect key ranges in serializable transactions.
PAGE	An 8-kilobyte (KB) page in a database, such as data or index pages.
EXTENT	A contiguous group of eight pages, such as data or index pages.
HoBT	A heap or B-tree. A lock protecting a B-tree (index) or the heap data pages in a table that does not have a clustered index.
TABLE	The entire table, including all data and indexes.
FILE	A database file.
APPLICATION	An application-specified resource.
METADATA	Metadata locks.
ALLOCATION_UNIT	An allocation unit.
DATABASE	The entire database.

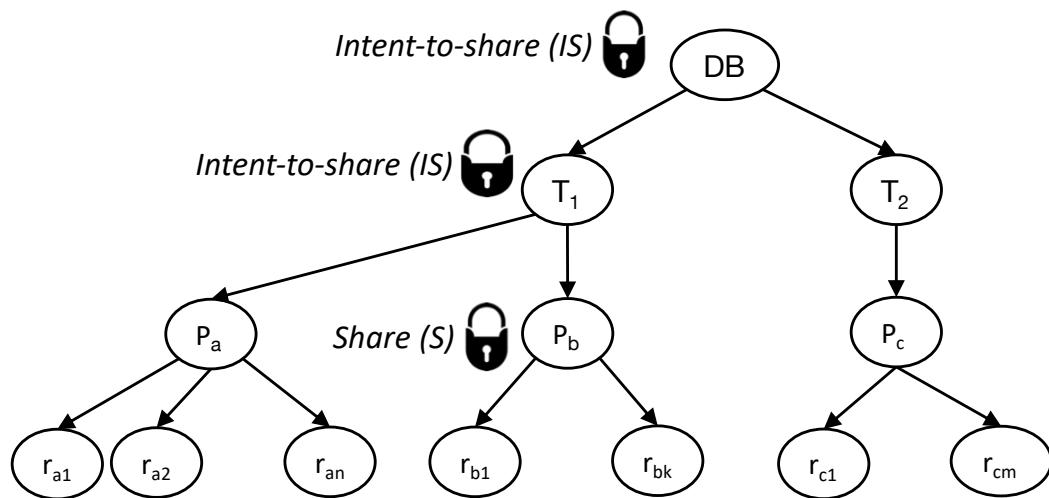
From MS SQL Server

[https://technet.microsoft.com/en-us/library/jj856598\(v=sql.110\).aspx](https://technet.microsoft.com/en-us/library/jj856598(v=sql.110).aspx)

New Lock Modes and Protocol



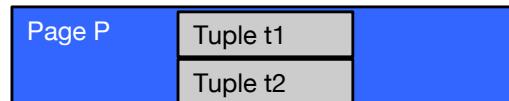
- Allow txns to lock at each level, but with a special protocol using new “intent” locks:
- Before getting S or X lock, txn must have proper intent locks on all its ancestors in the granularity hierarchy.



New Lock Modes – Intention Lock Modes



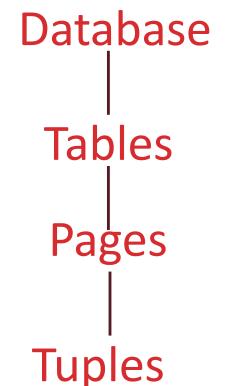
- 3 additional lock modes:
 - **IS:** Intent to get S lock(s) at finer granularity.
 - **IX:** Intent to get X lock(s) at finer granularity.
 - **SIX:** Like S & IX at the same time. Why useful?
- Intention locks allow a higher level node to be locked in S or X mode without having to check all descendent nodes



Multiple Granularity Locking Protocol



- Each txn starts from the root of the hierarchy.
- To get S or IS lock on a node, must hold IS or IX on parent node.
- To get X or IX or SIX on a node, must hold IX or SIX on parent node.
- Must release locks in bottom-up order.
- Enforce (strict) 2-phase locking as before
- Protocol is correct in that it is *equivalent to directly setting locks at leaf levels of the hierarchy.*
- What does the lock compatibility matrix look like?



Lock Compatibility Matrix

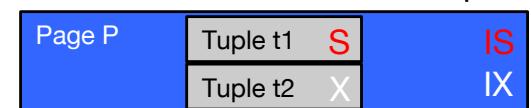
- IS – Intent to get S lock(s) at finer granularity.
- IX – Intent to get X lock(s) at finer granularity.
- SIX mode: Like S & IX at the same time.

	IS	IX	S	SIX	X
IS					
IX					
S			true		false
SIX					
X			false		false

Database
 |
 Tables
 |
 Pages
 |
 Tuples



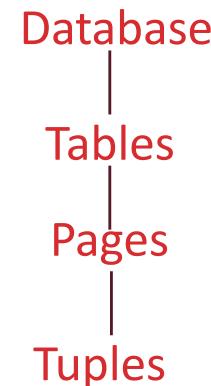
Handy simple case to remember:
Could 2 intent locks be compatible?



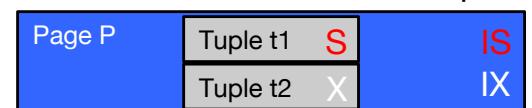
Lock Compatibility Matrix, Cont

- IS – Intent to get S lock(s) at finer granularity.
- IX – Intent to get X lock(s) at finer granularity.
- SIX mode: Like S & IX at the same time.

	IS	IX	S	SIX	X
IS	true	true	true	true	false
IX	true	true	false	false	false
S	true	false	true	false	false
SIX	true	false	false	false	false
X	false	false	false	false	false



Handy simple case to remember:
Could 2 intent locks be compatible?



Real-World Lock Compatibility Matrix



	NL	SCH-S	SCH-M	S	U	X	IS	IU	IX	SIU	SIX	UIX	BU	RS-S	RS-U	RI-N	RI-S	RI-U	RI-X	RX-S	RX-U	RX-X
NL	N	N	N	N	N	N	N	N	N	N	N	N	N	N	N	N	N	N	N	N	N	N
SCH-S	N	N	C	N	N	N	N	N	N	N	N	N	I	I	I	I	I	I	I	I	I	I
SCH-M	N	C	C	C	C	C	C	C	C	C	C	C	I	I	I	I	I	I	I	I	I	I
S	N	N	C	N	N	C	N	N	C	N	C	C	C	N	N	N	N	C	N	N	C	
U	N	N	C	N	C	C	N	C	C	C	C	C	C	N	C	N	N	C	C	N	C	C
X	N	N	C	C	C	C	C	C	C	C	C	C	C	C	C	N	C	C	C	C	C	C
IS	N	N	C	N	N	C	N	N	N	N	N	N	C	I	I	I	I	I	I	I	I	I
IU	N	N	C	N	C	C	N	N	N	N	N	N	C	I	I	I	I	I	I	I	I	I
IX	N	N	C	C	C	C	N	N	N	C	C	C	C	I	I	I	I	I	I	I	I	I
SIU	N	N	C	N	C	C	N	N	C	N	C	C	C	I	I	I	I	I	I	I	I	I
SIX	N	N	C	C	C	C	N	N	C	C	C	C	C	I	I	I	I	I	I	I	I	I
UIX	N	N	C	C	C	C	N	C	C	C	C	C	C	I	I	I	I	I	I	I	I	I
BU	N	N	C	C	C	C	C	C	C	C	C	C	C	C	N	I	I	I	I	I	I	I
RS-S	N	I	I	N	N	C	I	I	I	I	I	I	I	N	N	C	C	C	C	C	C	C
RS-U	N	I	I	N	C	C	I	I	I	I	I	I	I	N	C	C	C	C	C	C	C	C
RI-N	N	I	I	N	N	N	I	I	I	I	I	I	I	C	C	N	N	N	N	C	C	C
RI-S	N	I	I	N	N	C	I	I	I	I	I	I	I	C	C	N	N	N	C	C	C	C
RI-U	N	I	I	N	C	C	I	I	I	I	I	I	I	C	C	N	N	C	C	C	C	C
RI-X	N	I	I	C	C	C	I	I	I	I	I	I	I	C	C	N	C	C	C	C	C	C
RX-S	N	I	I	N	N	C	I	I	I	I	I	I	I	C	C	C	C	C	C	C	C	C
RX-U	N	I	I	N	C	C	I	I	I	I	I	I	I	C	C	C	C	C	C	C	C	C
RX-X	N	I	I	C	C	C	I	I	I	I	I	I	I	C	C	C	C	C	C	C	C	C

Key

N	No Conflict	SIU	Share with Intent Update
I	Illegal	SIX	Shared with Intent Exclusive
C	Conflict	UIX	Update with Intent Exclusive
NL	No Lock	BU	Bulk Update
SCH-S	Schema Stability Locks	RS-S	Shared Range-Shared
SCH-M	Schema Modification Locks	RS-U	Shared Range-Update
S	Shared	RI-N	Insert Range-Null
U	Update	RI-S	Insert Range-Shared
X	Exclusive	RI-U	Insert Range-Update
IS	Intent Shared	RI-X	Insert Range-Exclusive
IU	Intent Update	RX-S	Exclusive Range-Shared
IX	Intent Exclusive	RX-U	Exclusive Range-Update
		RX-X	Exclusive Range-Exclusive

From MS SQL Server

[https://technet.microsoft.com/en-us/library/jj856598\(v=sql.110\).aspx](https://technet.microsoft.com/en-us/library/jj856598(v=sql.110).aspx)

Phantom Problem



- So far we have assumed the database to be a *static* collection of elements (=tuples)
- If tuples are inserted/deleted then the *phantom problem* appears

Suppose there are two blue products, A1, A2:

Phantom Problem



T1

T2

```
SELECT *
FROM Product
WHERE color='blue'
```

```
INSERT INTO Product(name, color)
VALUES ('A3','blue')
```

```
SELECT *
FROM Product
WHERE color='blue'
```

Is this schedule serializable ?

Suppose there are two blue products, A1, A2:

Phantom Problem



T1

T2

```
SELECT *
FROM Product
WHERE color='blue'
```

```
INSERT INTO Product(name, color)
VALUES ('A3','blue')
```

```
SELECT *
FROM Product
WHERE color='blue'
```

Is this schedule serializable ?

No: T1 sees a “phantom” product A3

Suppose there are two blue products, A1, A2:

Phantom Problem



T1

```
SELECT *  
FROM Product  
WHERE color='blue'
```

T2

```
INSERT INTO Product(name, color)  
VALUES ('A3','blue')
```

```
SELECT *  
FROM Product  
WHERE color='blue'
```

But this is conflict-serializable!

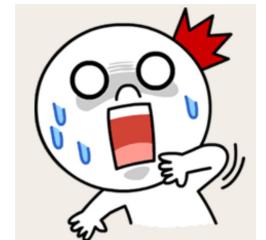
R₁(A1);R₁(A2);W₂(A3);R₁(A1);R₁(A2);R₁(A3)

W₂(A3);R₁(A1);R₁(A2);R₁(A1);R₁(A2);R₁(A3)

Phantom Problem



- A “phantom” is a tuple that is invisible during **part** of a transaction execution but not invisible during the **entire** execution
- In our example:
 - T1: reads list of products
 - T2: inserts a new product
 - T1: re-reads: a new product appears !
- Conflict-serializability assumes DB is *static*
- When DB is *dynamic* then c-s is not serializable.



Dealing With Phantoms



- Lock the entire table
- Lock the index entry for ‘blue’
 - If index is available
- Or use predicate locks
 - A lock on an arbitrary predicate

Dealing with phantoms is expensive !

Summary of Serializability



- Serializable schedule = equivalent to a serial schedule
- (strict) 2PL guarantees *conflict serializability*
 - What is the difference?
- **Static database:**
 - *Conflict serializability* implies serializability
- **Dynamic database:**
 - This no longer holds

Summary, cont.



- **Correctness criterion for isolation is “serializability”.**
 - In practice, we use “conflict serializability” which is conservative but easy to enforce
- **Two Phase Locking and Strict 2PL: Locks implement the notions of conflict directly**
 - The lock manager keeps track of the locks issued.
 - **Deadlocks** may arise; can either be prevented or detected.
- **Multi-Granularity Locking:**
 - Allows flexible tradeoff between lock “scope” in DB, and # of lock entries in lock table
- **More to the story**
 - Optimistic/Multi-version/Timestamp CC
 - Index “latching”, phantoms
 - Actually, there’s much much more ☺